

# Manual scheduling

Every pod has a field nodeName that is not default set

## How scheduling works

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
      - containerPort: 8080
  nodeName:
```

that, by default, is not set

5/1/2017

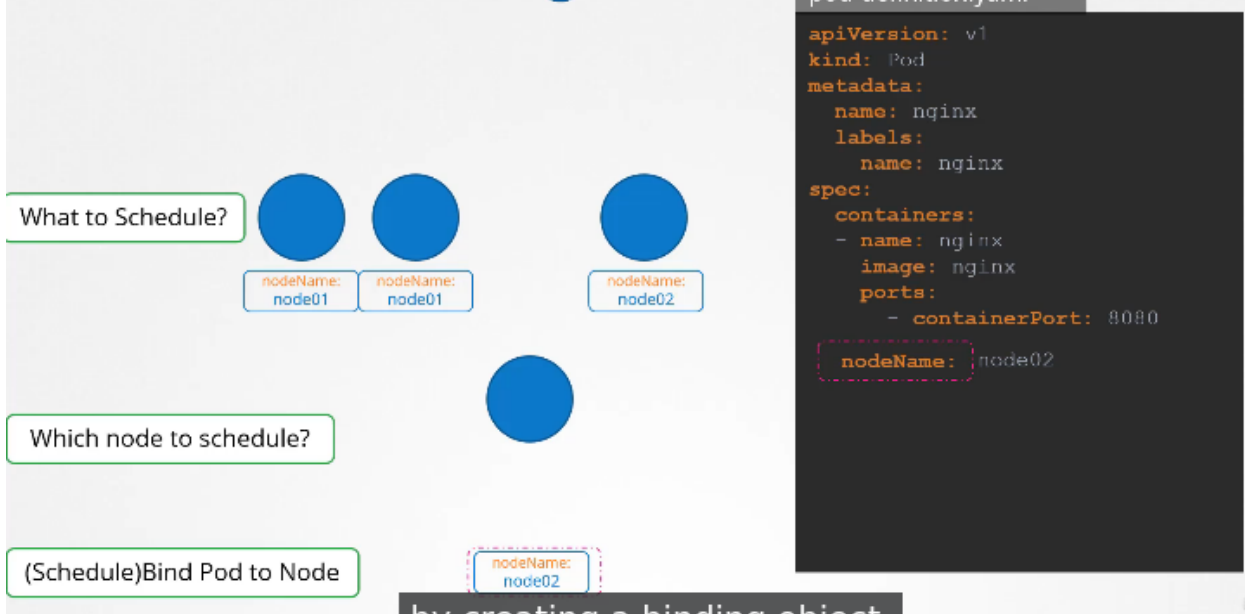
Scheduler goes through all pods and scan all pods which not have  
Thewse field that is nodeName  
And these are the candidate for scheduling

Then it find the right node for the pod ny running scheduling algoritrh

Once identified it scheduled the pod on the node  
By setting the node name property to the name of the node

By creating a binding object

# | How scheduling works



So if there is no scheduler to monitor and schedule node  
Then pod contains to be in a pending state

**Then in this case what we can do**

We can manually assign pods to nodes yourself

Without the scheduler the easiest way to schedule the pod is to set the  
Node name to the name of the node

# No Scheduler!

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Pending	0	3s

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
nginx	1/1	Running	0	9s	10.40.0.4	node02

```
pod-definition.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 8080
  nodeName: node02
```

in your pod specification file while creating the pod.

So that pod is assigned to specified node

## If pod is already created

K8s does not allow to modify nodeName property so how to assign node to a existing pod

## Create a binding object and send a POST request to the pod's binding API

Thus in that way what the actual scheduler does

In binding object we specified a target node as node of the name

And then send a POST request to the pod's binding API

With the data set to the binding object in a JSON format

```
apiVersion: v1
Kind: Binding
Metadata:
  Name: nginx
Target:
  apiVersion: v1
  Kind : Node
  Name : name-of-node
```

# No Scheduler!


## Pod-binding-definition.yaml

```
apiVersion: v1
kind: Binding
metadata:
  name: nginx
target:
  apiVersion: v1
  kind: Node
  name: node02
```

## pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
      - containerPort: 8080
```

```
curl --header "Content-Type:application/json" --request POST --data
http://$SERVER/api/v1/namespaces/default/pods/$PODNAME/binding/
```



Send a POST request to POD binding API

```
curl --header "Content-Type:application/json" --request POST --data '{"apiVersion":"v1", "kind": "Binding" ... }'
http://$SERVER/api/v1/namespaces/default/pods/$PODNAME/binding/
```

Curl -header "Content-Type:application/json" -request POST -data '{json format of binding object}' [https://\\$SERVER/api/v1/namespaces/default/pods/4PODNAME/binding/](https://$SERVER/api/v1/namespaces/default/pods/4PODNAME/binding/)

**Kubectl replace -force -f yaml-file.yaml**

We can't move pod from one node to another or system to system we  
First delete the pod and then move the pod

# Labels and Selector

Labels means group things together and selector filter things sbased on our needs  
Lablesla == properties attached to each item

Selector filter thse items based on label  
Every object has level

Objec t has levels  
And selector fuilter these object based on labels

Selector == condition

Labels:  
Key1: value1

## 2 ways

1 . kubectl get pods **--selector app=App1**

2 . in replica set

Selector:  
    matchLables  
        Key1: value1

We can specify many labels

On match rs created successfully

## 3 . for service

Selector:  
    Key1: value1

To group and select objects == labels and selector

# Annotations

It is used to record other details for informatory purpose

Like name , version , contact details . phone numbers , email ids etc

Buildversion

That may be used for some kind of integration purpose

## Annotations

```
replicaset-definition.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: simple-webapp
  labels:
    app: App1
    function: Front-end
  annotations:
    buildversion: 1.34
spec:
  replicas: 3
  selector:
    matchLabels:
      app: App1
  template:
    metadata:
      labels:
        app: App1
        function: Front-end
    spec:
      containers:
        - name: simple-webapp
          image: simple-webapp
```

Well, that's it for this lecture

Udemy

metadata:

Name:

labels:

annotations:

Buildversion: 1.34

**Kubectl get pods --selector env=dev --no-headers | wc -l**

**Kubectl get all --selector env=dev --no-headers | wc -l**

**Kubectl get pods --selector env=dev,bu=finance --no-headers | wc -l**

# Taints and Tolerations

Pod to node relationship


How we can restrict what pods are placed on what nodes

Ex:

Bug approaching a person

To prevent person from bug we need to spray the person with a repellent spray

Or taint



or a taint as we will call it in this lecture.

So that bug is intolerant to the smell so when bug approaches the person the taint throws the bug away

Bug == 1 intolerant to the smell or taint  
Or may be  
2 tolerant to the smell

But other bugs which are tolerant to this smell and so that bugs land on the person because taint do not do nothing



Bug land on the person or not

Taint on the person  
And toleration level of the bug


### In k8s

Person === node means we apply taint on the node

Bug === pod and we apply toleration on the pod



A B C D



the person is a node and the bugs are pods.

It does not lead any security or intrusion on the cluster

Taints and tolerations are used to set restrictions on what pods can be scheduled on a pod

Pod == a

When pods are created k8s scheduler tries to place these pod on the avail node

As when there is no restriction then scheduler places the pods across all the nodes to balancer the nodes equally



across all of the nodes to balance them out equally.

**nscript**

---

But let us assume that there is a particular node that have dedicated resources for a particular application or use case then in this case how we can deploy

So we want to place only those pods on these nodes that belongs these applications

So we can do

1, first prevent all pods from being placed on the node by placing a taint on the node  
Means apply taint on the node

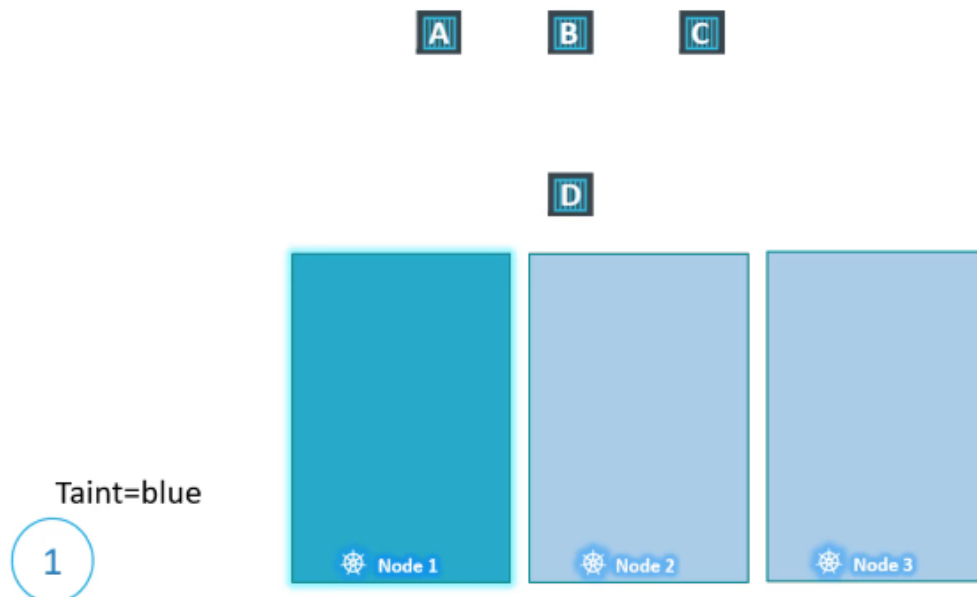
**Taint and tolerant == field or property of object**

Taint=blue

And by default pods have no tolerance means 0 tolerant level

Means none of the pods can tolerate any taint

So none of the pods can be placed on the taint node without having toleration power



on node one, as none of them can tolerate the taint blue.

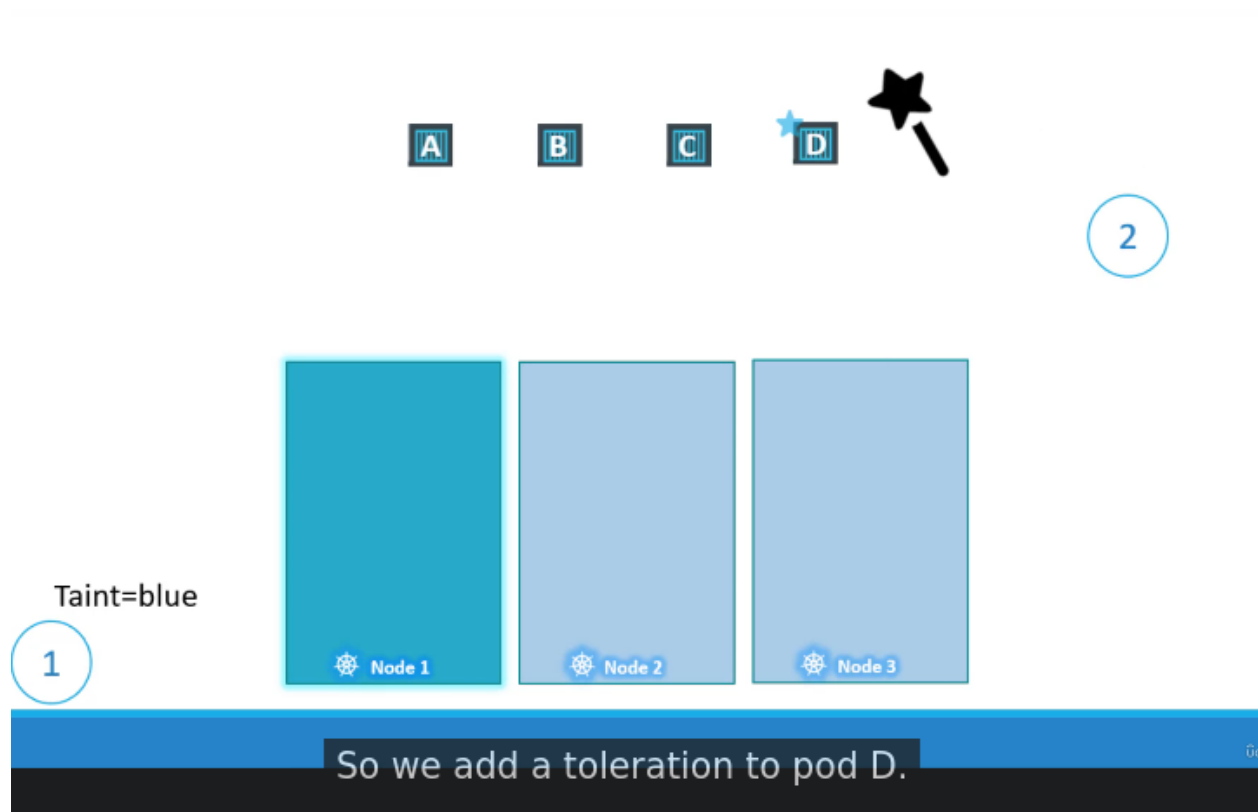
So to place node on the taint node we need to enable certain pods

To be placed on this node

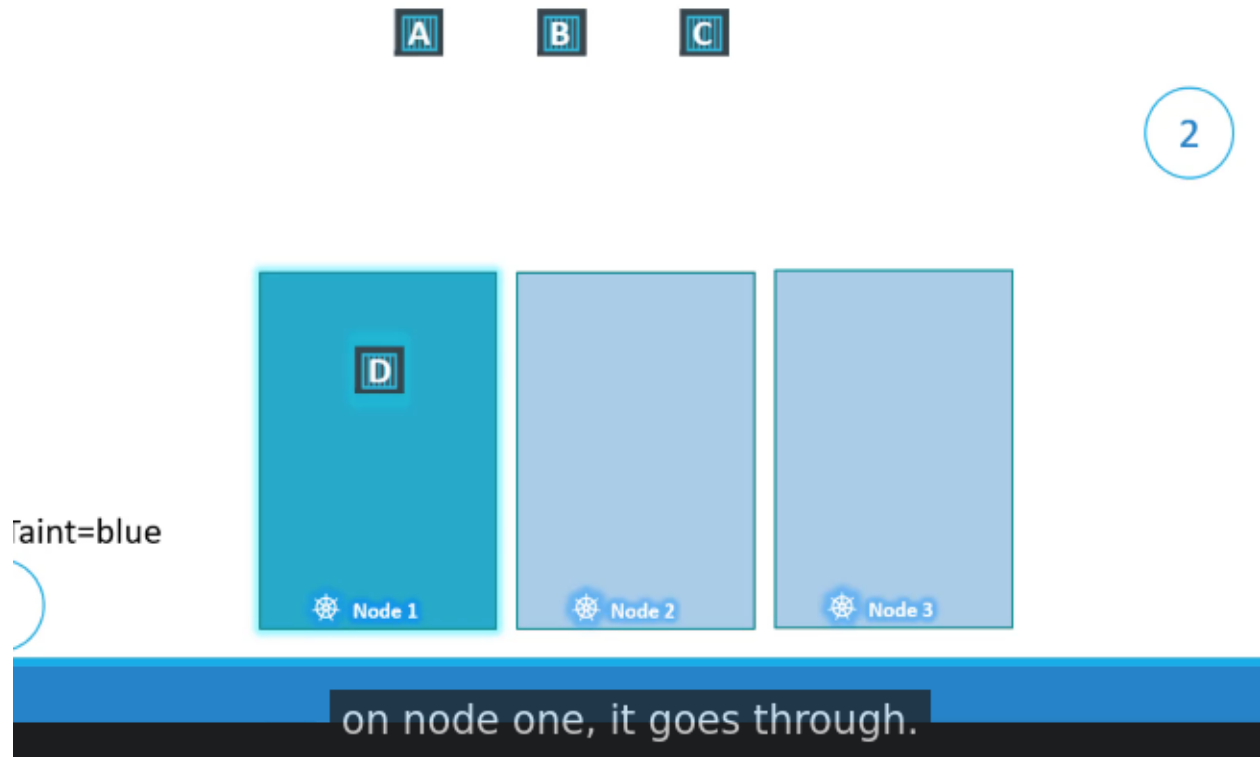
For this we need to specify which pods are tolerate to this taint node

Pod = a,b,c,d and we want only pod D to placed on the taint node

SO we need to add a toleration power to the pod D so that pod D is tolerate the node node 1



And so when scheduler tries to put pod D on node 1 it will now placed on node 1

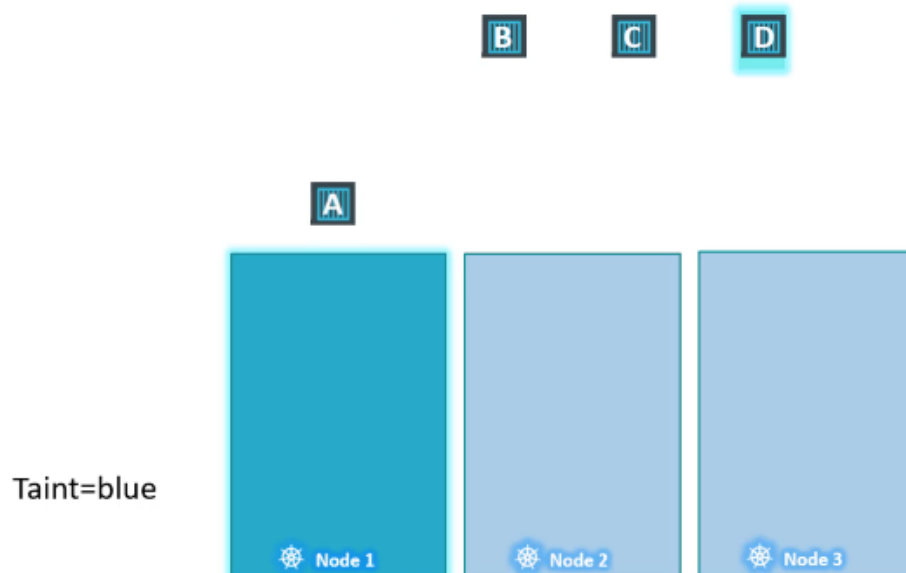


Node 1 only accept those pods which can tolerate taint blue


1 . scheduler tries to put pod A on node 1 but due to taint it will not placed  
And thrown off and then pod A goes to node 2 and node 2 has no taint then pod A will placed on node 2

Same with pod B node 1 thrown off pod B and then it will go to node 3

Same with C

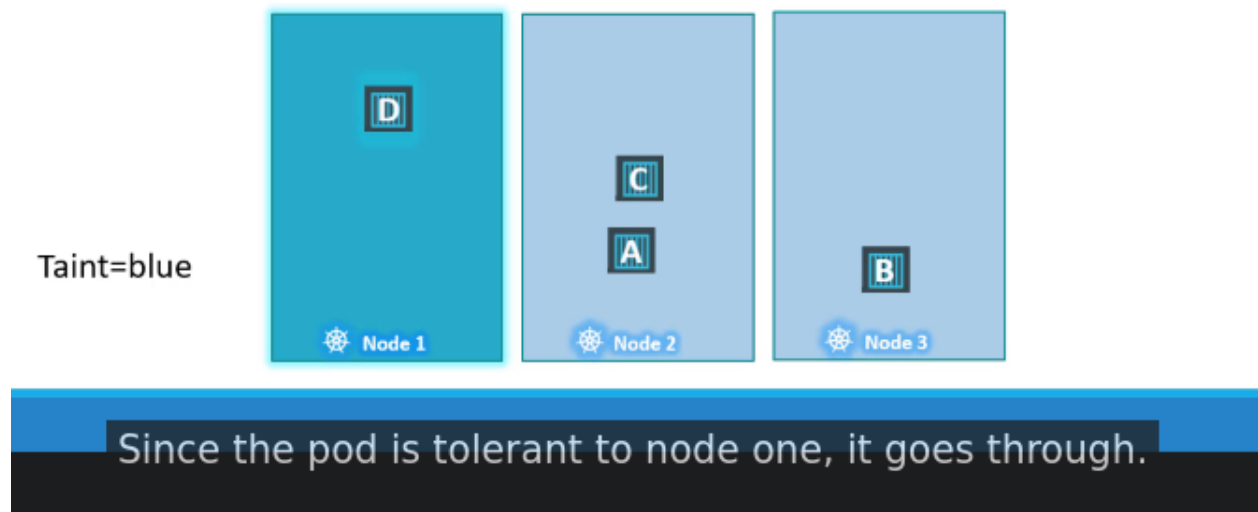


but due to the taint it is thrown off



It is thrown off again and ends up on node two.

2 . now pod D comes since the pod is tolerate to the taint node 1 , scheduler put the pod on node 1



**Taint = node**

**Toleration = pod**

Pod which have power to tolerate taint on taint node can be placed on that node

By default pod have no toleration

## Create

**KubectI taint node node-name key=value:taint-effect**

**Taint-effect**

**The** taint-effect defines what would happen if a pod do not tolerate the taint

There are 3 taint-effect

1. NoSchedule :

Means pod will not be scheduled on the node

2 PreferNoSchedule == means system will try to avoid a pod on the node but

That is not guaranteed

3 . NoExecute ==

Means that new pods will not be scheduled on the node and existing pods on the node , if any will be evicted if they do not tolerate the taint

नए पॉड्स को नोड पर शेड्यूल नहीं किया जाएगा और नोड पर मौजूदा पॉड्स, यदि कोई हो तो बेदखल कर दिया जाएगा

नए पॉड्स को नोड पर शेड्यूल नहीं किया जाएगा और नोड पर मौजूदा पॉड्स, अगर कोई दागी को बर्दाश्त नहीं करता है तो उसे बेदखल कर दिया जाएगा

These pods may have been scheduled on the node before the taint was applied to the node.

ये पॉड्स रहे होंगे

दागी को नोड पर लागू करने से पहले नोड पर निर्धारित किया गया था।

## Taints - Node

```
kubectl taint nodes node-name key=value:taint-effect
```

NoSchedule | PreferNoSchedule | NoExecute

What happens to PODs  
that do not tolerate this taint?

```
kubectl taint nodes node1 app=blue:NoSchedule
```

and an effect of no schedule.

Udemy



# Kubectll craete taint node1 app=blue:NoSchedule

## TOLERATIONS ARE ADDED TO THE PODS

### How to add toleration to a pod

First open pod definiton file

And under spec: section of pod-definiton file

Add a sectiions called **tolerations**

**And**

Move the same values while craeting the taint in double quotes

Spec:

Tolerations:

- Key: "app"  
Operator: "Equal"  
Value: "blue"  
Effect: "NoSchedule"

So now when we craete pod or updated with the new tolerations

Two things may happen

- 1 . Not scheduled on nodes ( if first time craeting:
- 2 . evicting from the existing nodes .

## ✦ Tolerations - PODs

---

```
kubectl taint nodes node1 :
```

```
pod-definition.yml
apiVersion:
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: nginx-container
      image: nginx
  tolerations:
    - key: "app"
      operator: "Equal"
      value: "blue"
      effect: "NoSchedule"
```

the no-execute taint effect in a bit more depth

0den

### NoExecute taint-effect

# Taint - NoExecute

---



In this example, we have three nodes running some workload.

**anscript**

**We do not have any taints and tolerations at this point**

Now we taint node1 and tolerate some pods which belong to node 1

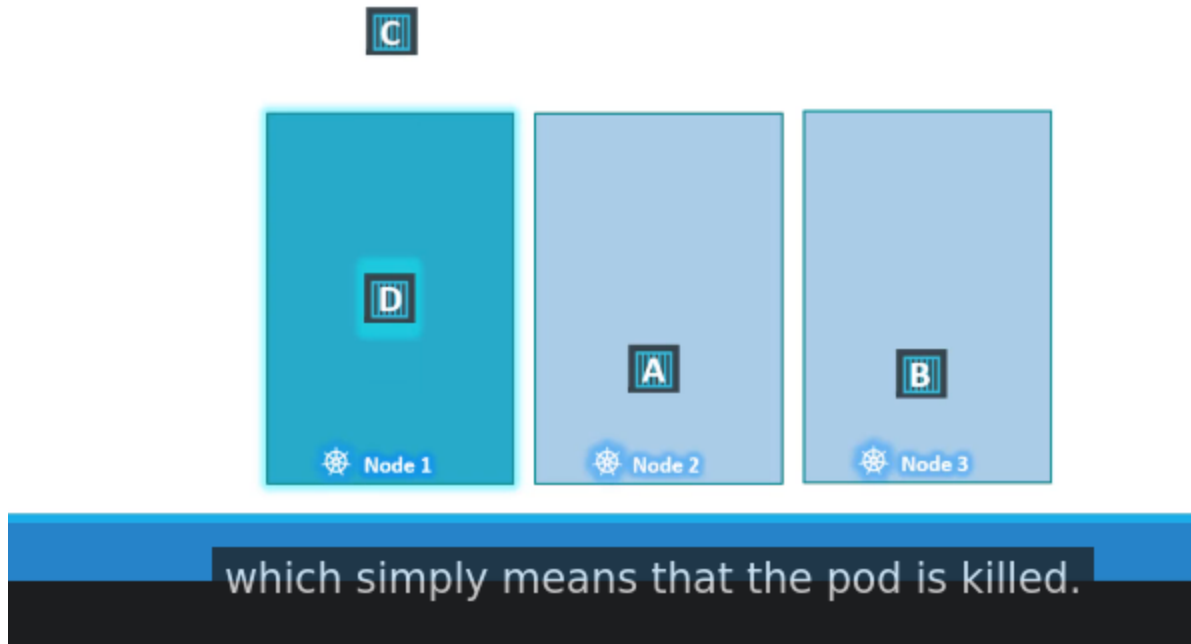
D = tolerate

And tainti-effect = NoExecute

And when applied pod c is evicted fro node 1 and it simply means kill

# Taint - NoExecute

---



Taints and tolerations are meant to specify nodes to accept certain pods only

In this case node 1 only accept pod 1

But tolerate pod can be placed on the other node also

**It only allows node to accept certain types of pods**

But if we want to place pod on only specific node we can achieve through node affinity

## Master nodes on the cluster

**It is** just another node

That has all the capabilities of hosting a pod and runs all the management software

Scheduler does not schedule any pods on the master node

When the k8s cluster is first set up

A taint is set on the master node automatically that prevents any pods from being scheduled on this node

We can see this and also modify this behaviour

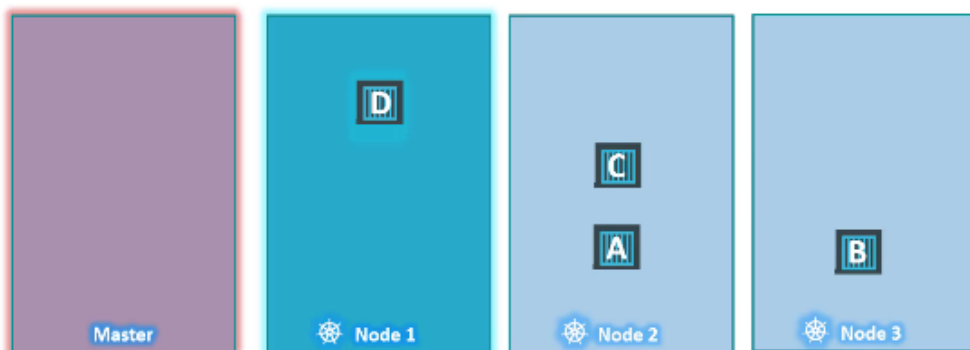
But best practice is to not deploy application workloads on a master server

**Kubectl describe node kubemaster | grep Taint**

**To delete taint**

**kubectl taint node controlplane node-role.kubernetes.io/control-plane:NoSchedule-**

```
kubectl describe node kubemaster | grep Taint
Taints:          node-role.kubernetes.io/master:NoSchedule
```



not schedule any pods on the master node.

**Nodes from accessing certain pods**

**Node 1 only accept pod D**

**But it does not guarantee pod D will always be placed on node one**

**If we want that pod should go to a certain pod not any random pod == we can achieve through node affinity**

## Kubectl taint node node-name key=value:NoScheduler

```
RestartPolicy: Always
status: {}

controlplane ~ → kubectl run bee --image=nginx --dry-run=client -o yaml > bee.yaml
controlplane ~ →
```

See 'kubectl taint -h' for help and examples

```
controlplane ~ * kubectl taint node controlplane node-role.kubernetes.io/control-plane:NoSchedule-
node/controlplane untainted

controlplane ~ →
```

Use "kubectl options" for a list of global command-line options (applies to all commands).

```
controlplane ~ → kubectl taint nodes node01 spray=mortein:NoSchedule
node/node01 tainted
```

## Node Selector

3 node cluster and 2 are smaller nodes with lower hardware resources one is larger node  
With higher resources

And you have different kinds of workload sin your cluster


For ex ;

Data processing workloads that require higher horsepower to the larger node


Because that have the resources that will not go out of resources

In case the job demand extra resources

But by default setup any pod can go to any nodes



You have a three node cluster of which two are smaller nodes



may very well end up on nodes two or three,

In this case pod C which require large resources may be placed in node c which have less resources

So to avoid this we use node selector

We set limitation on the pod so that they only run on particular nodes

**Two ways to do**

1. Node Selector == SIMPLE AND easier method

For this we look into the pod definition file created earlier

User spec section specify the field nideSelector and attribute size as large

Spec:

Container:

nodeSelector:

Size: Large

But from where size Large comes and how k8s know which has Large Node

**Basically the key value size and large are label key value pair assigned to nodes**

Scheduler uses these labels

To match and identify the right node to place the pods on


**For this we first label our nodes to use this one before creating this pod**



# Node Selectors

---

```
pod-definition.yml
apiVersion:
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: data-processor
      image: data-processor
  nodeSelector:
    size: Large
```



So, let us go back and see how we can label the nodes.

How we can the labels the nodes

**Kubectll label nodes node-name label-key=label-value**

**Kubectll label nodes node-1 size=Large**

Now we can create pod

With node selector as size Large


,

# Label Nodes

---

```
▶ kubectl label nodes <node-name> <label-key>=<label-value>
```

```
▶ kubectl label nodes node-1 size=Large
```



Now that we have labeled the node,

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) **[Reviews](#)** [Learning tools](#)

Now run command

# Node Selector



```
pod-definition.yml
apiVersion:
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
  - name: data-processor
    image: data-processor
  nodeSelector:
    size: Large
```

```
kubectl create -f pod-definition.yml
```

it is placed on node one as desired.

0dem

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)


We use single label and selector to achieve our goal  
We can't use mix of LABELS

# Node Selector - Limitations

---



- Large OR Medium?
- NOT Small



You cannot achieve this using node selectors.