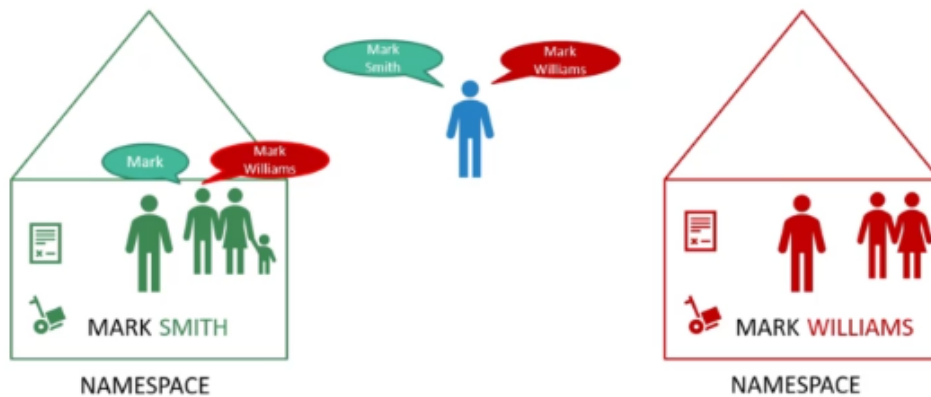


Namespaces

It is basically a environment

Set of rule and resources



These houses correspond to name spaces in Kubernetes.

Default namespace == it created auto by k8s when we create cluster

Set of pods and services for their internal such as networking solution
== kube-system == created at cluster startup

It is isolated from user so that user can't delete

Kube-public == it is public here resources made available to all users

Environment small == default


For enterprise == use your own created namespace

Dev and env == set of policies define who can do what

Resource quota to each namespaces

Resources within a namespace can refer by simply by their name

DNS



the web app pod can reach the DB service

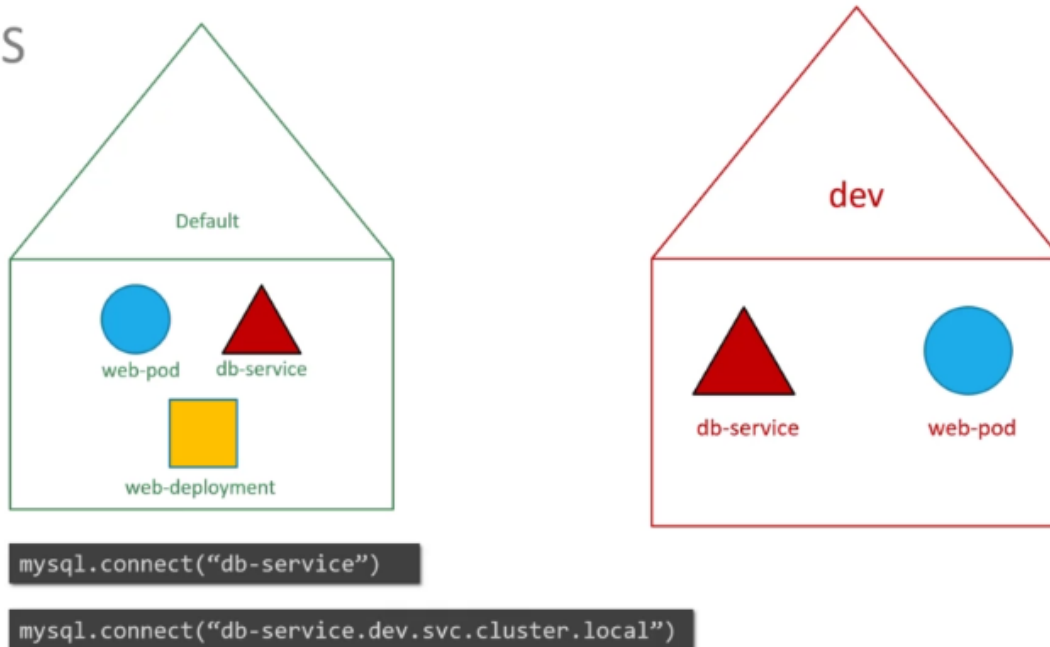
```
mysql.connect(db-serviuce)
```

For other namespace we used

Must append the name of namespace to the name of service

```
mysql.connect(db-service.dev.svc.cluster.local)
```

DNS



That would be dbservice.dev.svc.cluster.local.



[Course content](#)

[Overview](#)

[Q&A](#)

[Notes](#)

[Announcements](#)

[Reviews](#)

[Learning tools](#)

When a service is created a dns name is added automatically


Cluster.local == default domain name of the k8s cluster

Svc is the name of of sub domain od service

KubectI get pods == default

KubectI get pods --namespace=kube-system or -n

```
> kubectl create -f pod-definition.yml  
pod/myapp-pod created
```



```
pod-definition.yml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: myapp-pod  
  labels:  
    app: myapp  
    type: front-end  
spec:  
  containers:  
  - name: nginx-container  
    image: nginx
```

the pod is created in the default name space.

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)

To create in other space
Use namespace option

```
> kubectl create -f pod-definition.yml
```

```
pod/myapp-pod created
```

```
> kubectl create -f pod-definition.yml --namespace=dev
```

```
pod/myapp-pod created
```

```
pod-definition.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: myapp-pod
```

```
  labels:
```

```
    app: myapp
```


```
    type: front-end
```

```
spec:
```

```
  containers:
```

```
  - name: nginx-container
```

```
    image: nginx
```



Metadata:

Namespace: dev or prod

And then create pod

KubectI create -f yaml-file

```
> kubectl create -f pod-definition.yml
pod/myapp-pod created

> kubectl create -f pod-definition.yml
pod/myapp-pod created
```

```
pod-definition.yml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  namespace: dev
labels:
  app: myapp
  type: front-end
spec:
  containers:
  - name: nginx-container
    image: nginx
```

This is a good way to ensure your resources

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)

How to create namespace


```
apiVersion: v1
Kind: Namespace
Metadata:
  Name:dev
```

Kubectll create -f namespace-file.yaml

2 . 2ns way

Kubectll create namespace namespace-name

Create Namespace



```
namespace-dev.yml
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: dev
```

```
> kubectl create -f namespace-dev.yml
namespace/dev created
```

```
> kubectl create namespace dev
namespace/dev created
```

followed by the name of the name space.

Udemy

We have 3 ns dev , default , prod by default we are in default

How to switch

Kubectl get pods --namespace=dev

Kubectl get pods == output is default ns pod

Kubectl get pods --namespace=prod

But we dont want to mention ns always

Switch to another ns permanently

**Kubectl config set-context \$(kubectl config current-context)
--namespace=namespace-name**

And now only run

Kubectl get pods == output is only the dev environment

Pods in all namespaces


Kubectl get pods --all-namespaces

Kubectl get pods -A


Switch



```
> kubectl get pods --namespace=dev
```



```
> kubectl get pods
```



```
> kubectl get pods --namespace=prod
```

```
> kubectl config set-context $(kubectl config current-context) --namespace=dev
```

```
> kubectl get pods
```

```
> kubectl get pods --namespace=default
```

```
> kubectl get pods --namespace=prod
```

```
> kubectl config set-context $(kubectl config current-context) --namespace=prod
```

```
> kubectl get pods --namespace=dev
```

```
> kubectl get pods --namespace=default
```

```
> kubectl get pods
```

```
> kubectl get pods --all-namespaces
```

This will list all the pods in all of the name spaces.

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)

First find out current context and then add the namespace for that context

Limit resource

apiVersion: v1

Kind: ResourceQuota

Metadata:


Name: compute-quota

Namespace: dev

Spec: // provides your init
Hard:
 Pods: "10"
 Requests.cpu: "4"
 requests.memory: 5Gi
limits.cpu: "10"
Limits.memory: 10Gi

Kubectl create -f resource-quota-file-name.yaml

Resource Quota



```
Compute-quota.yaml
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-quota
  namespace: dev
spec:
  hard:
    pods: "10"
    requests.cpu: "4"
    requests.memory: 5Gi
    limits.cpu: "10"
    limits.memory: 10Gi
```

```
> kubectl create -f compute-quota.yaml
```

10 GB byte of memory, etcetera.

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)

Kubectl get ns or namespaces

Kubectl get ns or namespaces

Kubectl get pods – namespace=namespace-name

Kubectl get pods – n=namespace-name

Imperative and Declarative

Steps by step ---- final destination

How to do what to do ----- what to do

In iaac == a ser of instruction -== command ==Imperatibve

In iaac declare only requirements == Declarative === all things are done by software

There are 7 steps

In first run only 4 step execute

Then in next run

We need to provide checks that if this happen then dont apply

Im

Imperative

Kubernetes

Imperative

```
> kubectl run --image=nginx nginx
> kubectl create deployment --image=nginx nginx
> kubectl expose deployment nginx --port 80
> kubectl edit deployment nginx
> kubectl scale deployment nginx --replicas=5
> kubectl set image deployment nginx nginx=nginx:1.18
```

We can also use file

Create , replace delete

```
> kubectl set image deployment nginx nginx=nginx:1.18
```

```
> kubectl create -f nginx.yaml
```

```
> kubectl replace -f nginx.yaml
```

```
> kubectl delete -f nginx.yaml
```

Declarative

And deleting an object using the kubectl delete command.

Declarative

Declarative

```
> kubectl apply -f nginx.yaml
```

for creating, updating, or deleting an object.

[Course content](#)

[Overview](#)

[Q&A](#)

[Notes](#)

[Announcements](#)

[Reviews](#)

[Learning tools](#)

Apply command look existing configuration and figure out what changes need to be done to the system

Imperative commands

Imperative Commands

Create Objects

```
> kubectl run --image=nginx nginx
```

```
> kubectl create deployment --image=nginx nginx
```

```
> kubectl expose deployment nginx --port 80
```

Update Objects

```
> kubectl edit deployment nginx
```

```
> kubectl scale deployment nginx --replicas=5
```

```
> kubectl set image deployment nginx nginx=nginx:1.18
```

such as the run, create, or expose commands

Also edit scale set commands to update existing object

Run once and available only in session history

Yaml file

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

nginx.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```



Local file

pod-definition

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx:1.18
status:
  conditions:
    - lastProbeTime: null
      status: "True"
      type: Initialized
```



Kubernetes Memory

you're only left with your local definition file,

Odem

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

```
nginx.yaml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
  - name: nginx-container
    image: nginx
```

which in fact has the old image name in it.

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)

But this will not change local file

So to change in local dictl change and run following command

Kubectl replace -f yaml-file

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

```
> kubectl replace -f nginx.yaml
```

```
nginx.yaml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
  - name: nginx-container
    image: nginx:1.18
```

This way, going forward, the changes made are recorded

Completely delete and recreate objects

Kubectl replace --force -f yaml-file

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

```
> kubectl replace -f nginx.yaml
```

```
> kubectl replace --force -f nginx.yaml
```

```
nginx.yaml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
  - name: nginx-container
    image: nginx:1.18
```

but with the force option, like this.

[Course content](#) [Overview](#) [Q&A](#) [Notes](#) [Announcements](#) [Reviews](#) [Learning tools](#)

In imperative approach

If we create object and it already exists it gives error

When we run replace command object must be available

So to overcome all this problem we use declarative approach

Declarative

Declarative

Create Objects

```
> kubectl apply -f nginx.yaml
```

```
> kubectl apply -f /path/to/config-files
```

Update Objects

```
nginx.yaml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
  - name: nginx-container
    image: nginx
```

That way, all the objects are created at once.

And now when changes need to be made we need to change only local file and run apply command again

Apply command knows that object exists then he only update the objects with new changes

Exam Tips

Create Objects

```
> kubectl apply -f nginx.yaml
```

```
> kubectl run --image=nginx nginx
```

```
> kubectl create deployment --image=nginx nginx
```

```
> kubectl expose deployment nginx --port 80
```

Update Objects

```
> kubectl apply -f nginx.yaml
```

```
> kubectl edit deployment nginx
```

```
> kubectl scale deployment nginx --replicas=5
```

```
> kubectl set image deployment nginx nginx=nginx:1.18
```

you could use the imperative approach

Udemy

Certification Tips - Imperative Commands with Kubectl

While you would be working mostly the declarative way - using definition files, imperative commands can help in getting one time tasks done quickly, as well as generate a definition template easily. This would help save considerable amount of time during your exams.

Before we begin, familiarize with the two options that can come in handy while working with the below commands:

--dry-run: By default as soon as the command is run, the resource will be created. If you simply want to test your command, use the --dry-run=client option. This will not create the resource, instead, tell you whether the resource can be created and if your command is right.

-o yaml: This will output the resource definition in YAML format on screen.

Use the above two in combination to generate a resource definition file quickly, that you can then modify and create resources as required, instead of creating the files from scratch.

POD

Create an NGINX Pod

```
kubectl run nginx --image=nginx
```

Generate POD Manifest YAML file (-o yaml). Don't create it(--dry-run)

```
kubectl run nginx --image=nginx --dry-run=client -o yaml
```

Deployment

Create a deployment

```
kubectl create deployment --image=nginx nginx
```

Generate Deployment YAML file (-o yaml). Don't create it(--dry-run)

```
kubectl create deployment --image=nginx nginx --dry-run=client -o yaml
```

Generate Deployment with 4 Replicas

```
kubectl create deployment nginx --image=nginx --replicas=4
```

You can also scale a deployment using the kubectl scale command.

```
kubectl scale deployment nginx --replicas=4
```

Another way to do this is to save the YAML definition to a file and modify

```
kubectl create deployment nginx --image=nginx --dry-run=client -o yaml >  
nginx-deployment.yaml
```

You can then update the YAML file with the replicas or any other field before creating the deployment.

Service

Create a Service named redis-service of type ClusterIP to expose pod redis on port 6379

```
kubectl expose pod redis --port=6379 --name redis-service --dry-run=client -o yaml
```

(This will automatically use the pod's labels as selectors)

Or

`kubectl create service clusterip redis --tcp=6379:6379 --dry-run=client -o yaml` (This will not use the pods labels as selectors, instead it will assume selectors as **app=redis**. [You cannot pass in selectors as an option](#). So it does not work very well if your pod has a different label set. So generate the file and modify the selectors before creating the service)

Create a Service named nginx of type NodePort to expose pod nginx's port 80 on port 30080 on the nodes:

```
kubectl expose pod nginx --type=NodePort --port=80 --name=nginx-service --dry-run=client -o yaml
```

(This will automatically use the pod's labels as selectors, [but you cannot specify the node port](#). You have to generate a definition file and then add the node port in manually before creating the service with the pod.)

Or

```
kubectl create service nodeport nginx --tcp=80:80 --node-port=30080 --dry-run=client -o yaml
```

(This will not use the pods labels as selectors)

Both the above commands have their own challenges. While one of it cannot accept a selector the other cannot accept a node port. I would recommend going with the `kubectl expose` command. If you need to specify a node port, generate a definition file using the same command and manually input the nodeport before creating the service.

Reference:

<https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands>

<https://kubernetes.io/docs/reference/kubectl/conventions/>

`--labels="app=run"`

```
controlplane ~ + kubectl run redis --image=redis:alpine --labels="tier=db"
pod/redis created

controlplane ~ +
```

Create a service `redis-service` to expose the `redis` application within the cluster on port `6379`.

Use imperative commands.

Check

- Service: `redis-service`
- Port: `6379`
- Type: `ClusterIP`

```
service/redis-service exposed

controlplane ~ + kubectl get svc redis
Error from server (NotFound): services "redis" not found

controlplane ~ + kubectl get svc redis-service
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)
redis-service ClusterIP      10.43.56.187   <none>         6379/TCP

controlplane ~ + kubectl describe svc redis-service
Name:         redis-service
Namespace:    default
Labels:       tier=db
Annotations:   <none>
Selector:     tier=db
Type:         ClusterIP
IP Family Policy: SingleStack
IP Families:  IPv4
IP:           10.43.56.187
IPs:          10.43.56.187
Port:         <unset> 6379/TCP
TargetPort:   6379/TCP
Endpoints:    10.42.0.10:6379
Session Affinity: None
Events:       <none>

controlplane ~ +
```

1 2 3 4 5 6 7 8 9

Create a new pod called `custom-nginx` using the `nginx` image and expose it on `container` port `8080`.

Next

✓ Pod created correctly?

```
controlplane ~ + kubectl run custom-nginx --image=nginx --port=8080
pod/custom-nginx created

controlplane ~ +
```

Kubectl run httpd --image=httpd --port=80 --expose=true

Apply command

Local

last applied

live object